

# Advanced SQL & Query Optimization

Dr. Medha Kirti



## **Agenda**

- Nested Queries
- Correlated Subqueries
- Set Operations
  - Union
  - Intersect
  - Minus
- Query Optimization Basics
  - Indexes
  - Execution Plans

# Nested Queries

**Definition:** A query within another SQL query.

**Structure:**

SELECT ... FROM ... WHERE column = (SELECT ... FROM ...);

**Example:**

```
SELECT employee_id, name  
FROM employees  
WHERE department_id IN (SELECT department_id FROM departments  
WHERE location_id = 1000);
```

# Nested Queries

Employees:

EmpID	Name	Salary
1	Alice	5000
2	Bob	7000
3	Carol	6000

**Task:** Find employees who earn **more than the average salary**.

```
SELECT Name, Salary
FROM Employees
WHERE Salary > (
    SELECT AVG(Salary)
    FROM Employees
);
```

# Nested Queries

## How it Works:

1. Inner Query:

```
SELECT AVG(Salary) FROM Employees;
```

→ Calculates  
average salary  
(say it's 6000)

2. Outer Query:

```
SELECT Name, Salary FROM Employees WHERE Salary > 6000;
```

→ Returns  
employees with  
salary above 6000  
→ **Bob.**

## Output:

Name	Salary
Bob	7000

# Correlated Subqueries

**Definition:** A subquery that references columns from the outer query.

**Structure:**

```
SELECT ...  
FROM OuterTable AS o  
WHERE EXISTS (  
    SELECT ...  
    FROM InnerTable AS i  
    WHERE i.column = o.column  
);
```

# Correlated Subqueries

**Employees:**

EmpID	Name	DeptID	Salary
1	Alice	10	5000
2	Bob	10	7000
3	Carol	20	6000
4	Dave	20	5500

**Task:** Find employees who **earn more than the average salary of their department**.

```
SELECT Name, Salary, DeptID
FROM Employees AS E
WHERE Salary > (
    SELECT AVG(Salary)
    FROM Employees AS E2
    WHERE E2.DeptID = E.DeptID
);
```

# Correlated Subqueries

## How it Works:

- For **each employee**, the subquery calculates the **average salary** in their **own department**.
- Then compares their salary with that average.

## Output:

Name	Salary	DeptID
Bob	7000	10





Interview Question

## What is the difference between Normal Subquery and Correlated Subquery?

### Normal (Non-Correlated) Subquery

#### ◆ How It Works:

- The **inner query runs once**.
- The result is then used by the outer query.
- The inner query is **independent** of the outer query.

### Correlated Subquery

#### ◆ How It Works:

- The **inner query runs once for each row** of the outer query.
- The inner query **depends on the current row** of the outer query.
- Slower but more powerful and flexible.



Interview Question

**Describe a real-life scenario where you would use a correlated subquery instead of a JOIN.**

**Scenario:** Find employees earning more than the average of their own department.

**Why not JOIN?** Because we need the **average per department**, and compare it **row by row**.

**JOINS are better for combining related tables, but correlated subqueries are better for row-wise comparisons.**



Interview Question

Which of the following queries finds employees who earn more than the average salary in their own department? (Amazon)

- 1. A. Uses a simple GROUP BY
- B. Uses a correlated subquery**
- C. Uses a JOIN only
- D. Uses UNION

**Answer: B — Requires a correlated subquery comparing each employee to their department's average.**

# Set Operations

**Overview:** Operations that combine results from two or more queries.

## Types:

- Union
- Intersect
- Minus

## Important Rules:

- The **number and order of columns must match** in both SELECT queries.
- The **data types** of the columns must be compatible

# Union

**Definition:** Combines the results of two queries, removing duplicates.

**Example:**

```
SELECT name FROM employees  
UNION  
SELECT name FROM contractors;
```

# Intersect

**Definition:** Returns only the rows that are present in both queries.

**Example:**

```
SELECT name FROM employees  
INTERSECT  
SELECT name FROM contractors;
```

# Minus

**Definition:** Returns rows from the first query that are not in the second.

**Example:**

```
SELECT name FROM employees  
MINUS  
SELECT name FROM contractors;
```

# Quick Summary Table

<u>Operation</u>	<u>Purpose</u>	<u>Removes Duplicates?</u>	<u>Order of Columns Must Match?</u>
<b>UNION</b>	Combine and return distinct values	Yes	Yes
<b>UNION ALL</b>	Combine and return all values	No	Yes
<b>INTERSECT</b>	Return common records	Yes	Yes
<b>MINUS</b>	Return records from first query only	Yes	Yes





Interview Question

**Can we use INTERSECT and MINUS in all databases?**

**No, not all relational databases support INTERSECT and MINUS directly.**

Database System	INTERSECT	MINUS / EXCEPT	Notes
Oracle	✓ Supported	✓ MINUS is used	Fully supported
PostgreSQL	✓ Supported	✓ EXCEPT is used	Fully supported
SQL Server	✓ Supported	✓ EXCEPT is used	Fully supported
MySQL (before v8.0)	✗ Not Supported	✗ Not Supported	Use JOIN, IN, or NOT EXISTS instead
MySQL 8.0+	✗ Not Supported	✗ Not Supported	Still lacks direct support
SQLite	✓ Supported	✓ EXCEPT is used	Works with some limitations



Interview Question

**Explain a scenario where using INTERSECT is better than using INNER JOIN.**

**Use INTERSECT when you only need common values from identical columns in two tables without needing to combine additional columns. It is shorter and cleaner than a JOIN followed by DISTINCT.**

◆ Scenario:

Suppose you are analyzing customer behavior for an e-commerce platform (like Amazon or Flipkart). You want to find the `customer_ids` that appear in both the `Online_Customers` and `Store_Customers` tables — meaning customers who have shopped through both channels.

You do not need other details like names or addresses — just the matching IDs.

### Using INTERSECT:

```
SELECT customer_id FROM  
Online_Customers  
INTERSECT  
SELECT customer_id FROM  
Store_Customers;
```

- ✓ This gives you only the common customer\_ids.
- ✓ No duplicates (by default).
- ✓ Very readable and clear intent.

### Using INNER JOIN:

```
SELECT DISTINCT o.customer_id  
FROM Online_Customers o  
INNER JOIN Store_Customers s  
ON o.customer_id = s.customer_id;
```

- ✓ Same result, but requires:
  - Explicit aliasing
  - DISTINCT to remove duplicates (if needed)



Interview Question

What would happen if the columns selected in a set operation do not match in type or number?

The SQL engine will return an error. Set operations require the queries to return the **same number of columns** with **compatible data types**.

# Query Optimization Basics

**Importance:** Enhances performance and efficiency of SQL queries.

**Key Concepts:**

- Indexes
- Execution Plans

# Indexes

**Definition:** A database structure that improves the speed of data retrieval.

**Types:**

- Single-column Index
- Composite Index

```
CREATE INDEX idx_employee_name ON employees(name);
```

# Execution Plans

**Definition:** A detailed breakdown of how a SQL query will be executed.

**How to View:**

- Use 'EXPLAIN' keyword before your query.

**Example:**

```
EXPLAIN SELECT * FROM employees WHERE department_id = 10;
```



Interview Question

## What is the role of an index in query optimization?

An index speeds up data retrieval by allowing the database to avoid full table scans. It works like a lookup system and improves performance for SELECT queries.





Interview Question

**Describe a scenario at Amazon where query optimization is crucial.**

During product search, Amazon needs to retrieve millions of items fast. Indexes on product\_name, rating, and price allow optimized search filtering.

### **Why Query Optimization Matters**

- When a user searches or filters products (e.g., “wireless headphones under \$50 with 4+ star ratings”), Amazon’s database must:
- Quickly fetch relevant results
- Sort and filter efficiently
- Handle millions of simultaneous queries



Interview Question

Explain a situation where adding an index can decrease performance.

## 1. Frequent INSERT, UPDATE, DELETE operations

Why?

- Every time a row is **inserted, updated, or deleted**, the database must **update the index** as well.
- This **adds overhead** and slows down write operations.

## 2. Too Many Indexes on a Table

Why?

- Each index must be maintained separately.
- More indexes = more disk I/O = slower performance during **data modification**.

# Thank You