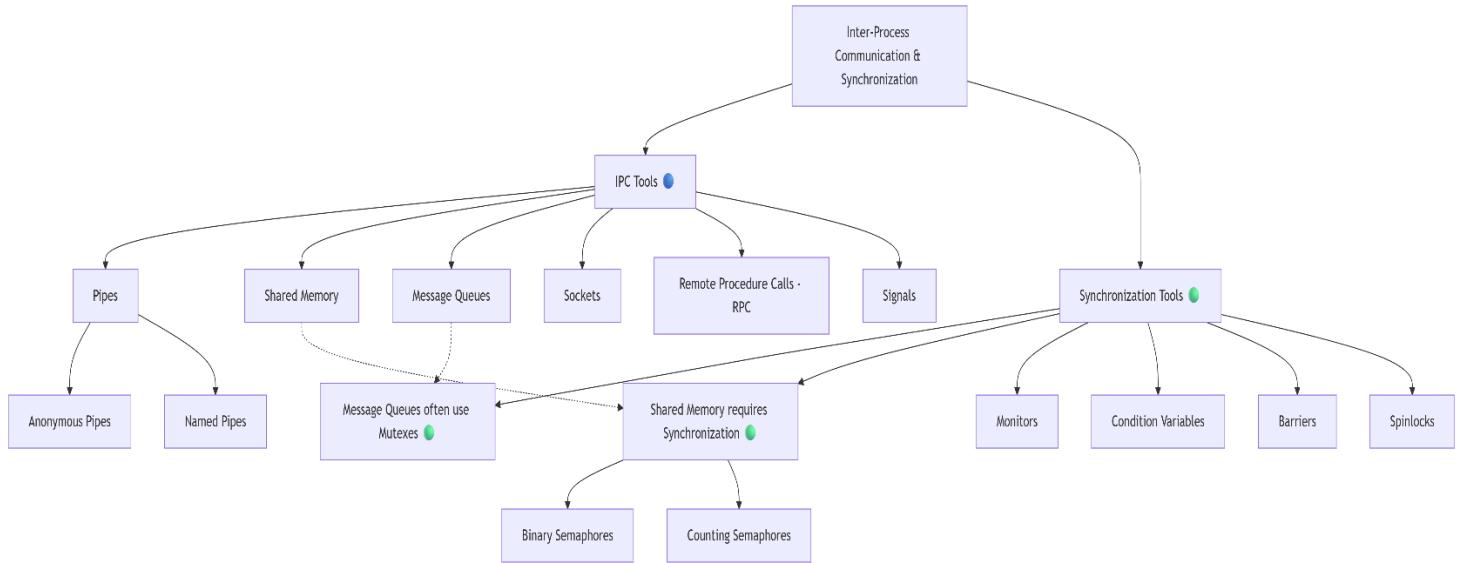


1.Slow Learner Activity:-



Question 1: IPC Focus

Q: Two processes need to exchange data on the same machine. One is a parent process, the other is its child. Which IPC tool should they use? Explain why.

Solution:

- **Tool:** Anonymous Pipes (●).
- **Why:**
 - Anonymous pipes are designed for one-way communication between parent-child processes.
 - They act like a temporary "data tunnel" (e.g., parent → child).
 - **Example:** In Linux, created via `pipe()` syscall before `fork()`.

Question 2: Synchronization Focus

Q: A critical resource (e.g., a printer) can only serve one process at a time. Which synchronization tool ensures exclusive access? How does it work?

Solution:

- **Tool:** Mutex (●).
- **How it works:**
 1. A process locks the mutex before using the resource.
 2. If another process tries to lock it, it waits (blocks).
 3. When done, the process unlocks the mutex.
- **Analogy:** Like a single key to a bathroom – only the key holder can enter.

Question 3: Combined IPC + Synchronization

Q: Two processes use Shared Memory (●) to edit the same file. Why might they corrupt data? Which synchronization tool (●) fixes this?

Solution:

- Problem:
Both processes might write to the same memory region simultaneously, causing a race condition (e.g., overwriting each other's data).
- Synchronization Tool: Mutex or Semaphore (●).
- How it helps:
 - Processes lock the mutex/semaphore before editing.
 - Only one process accesses the shared memory at a time.
 - Example:

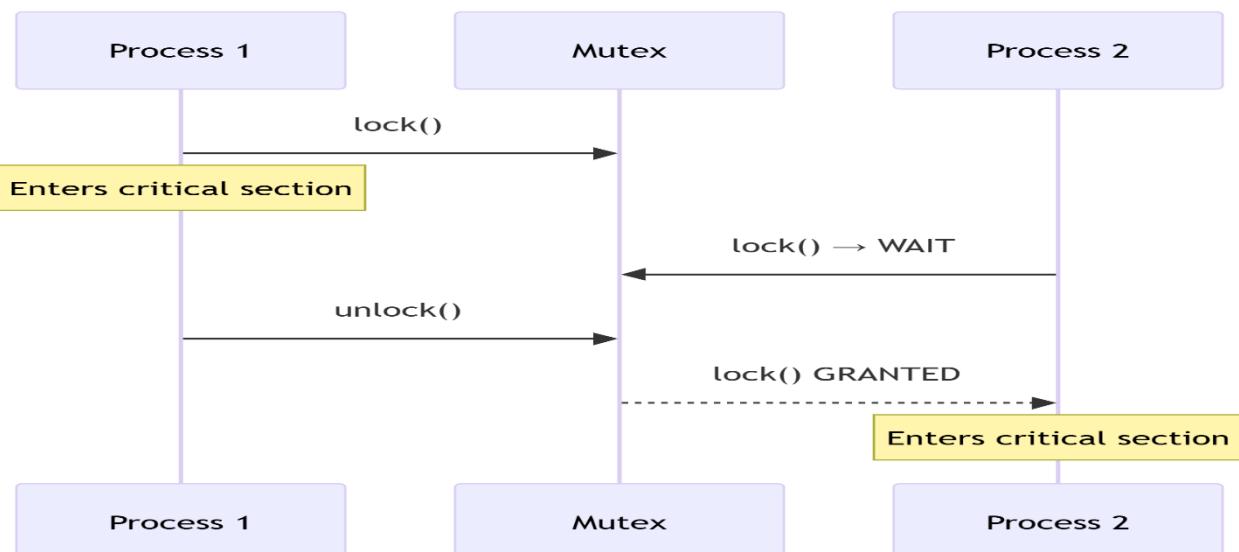
Code:

```
lock(mutex);      // Enter critical section
edit_shared_file();
unlock(mutex);    // Exit
```

Summary for Learners:

1. IPC Tools (●) = *Communication* (e.g., Pipes, Shared Memory).
2. Sync Tools (●) = *Coordination* (e.g., Mutex for exclusive access).
3. Shared Memory requires Sync – Always pair ● with ● to prevent chaos!

2.



Question 1: Mechanism

Q: Process A locks a mutex to access a shared printer. What happens when Process B attempts to lock the same mutex?

Solution:

- Process B **blocks (waits)** until Process A unlocks the mutex.
- *Why?* Mutexes enforce strict mutual exclusion – only one process can hold the lock at a time.

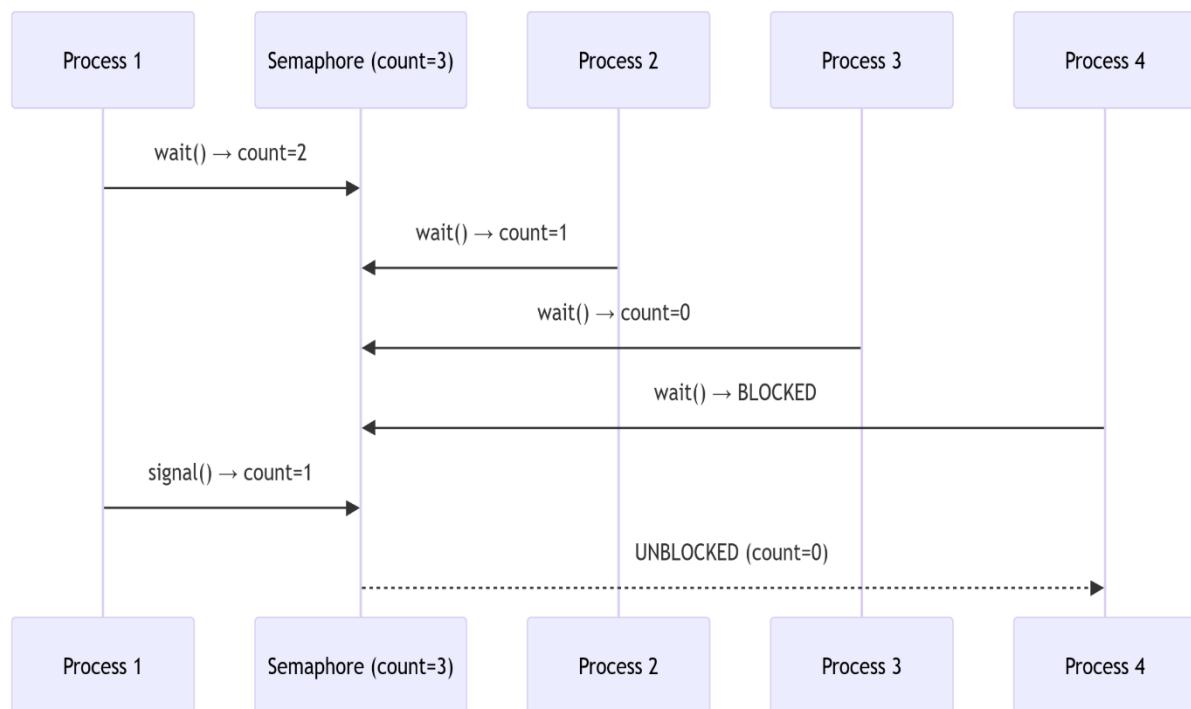
Question 2: Printer Use Case

Q: Why must a printer use a mutex instead of allowing concurrent access? What happens without it?

Solution:

- **Why?** Printers cannot handle mixed data from multiple processes (e.g., half Page 1 from Process A + half Page 1 from Process B = garbage).
- **Without Mutex:**
 - Processes print simultaneously → **corrupted output**.
 - Example: A letter and a report print blended into unreadable gibberish.

Moderate Learner Learner:--



3.

Question 1: Resource Allocation

Q: The semaphore starts with 3 resources. After Processes 1, 2, and 3 each acquire a resource via `wait()`, how many resources remain? What happens when Process 4 tries to acquire one?

Solution:

- **Resources left: 0** ($3 - 3 = 0$).
- **Process 4: Blocks immediately** because no resources are available.
- **Key concept:** `wait()` is a **resource request** – it fails if count=0.

Question 2: Signal Operation

Q: After Process 4 is blocked, Process 1 releases its resource with `signal()`. Why does the semaphore count temporarily become 1 but then drop to 0? What happens to Process 4?

Solution:

1. **Process 1's signal():**
 - Releases 1 resource → count increases from **0 → 1**.

2. Process 4:

- Is **unblocked automatically** and immediately consumes the resource.
- Performs `wait()` → count decreases from **1 → 0**.
- *Why?* `signal()` releases a resource AND wakes a waiting process.

3.

Which of the following are not shared by the threads of the same process?

- (a) Stack
- (b) Registers
- (c) Address Space
- (d) Message Queue

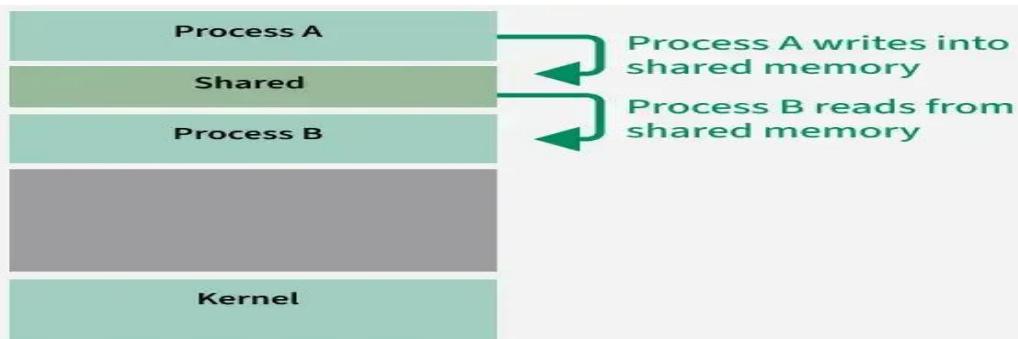
1. (a) and (d)
2. (b) and (c)
3. (a) and (b)
4. More than one of the above
5. None of the above

→ Solution:--

Option 3 : (a) and (b)

Operating Systems Question 3 Detailed Solution

- A thread is a lightweight process that can be managed independently by a scheduler.
- A message queue is a linked list of messages stored within the kernel and identified by a message queue identifier.
- Threads share the code and data segments and the heap, but they don't share the stack and registers. Message queue comes under data and therefore it is shared by thread.
- Threads of the same process don't share stack and register, hence option 3 is correct



1. In the shared memory diagram, if Process A writes data and Process B reads *at the same time* without synchronization, what problem occurs?

Answer:

Data corruption – Process B might read partially updated/wrong values.

Why? (Simple Explanation)

1. Write Operation isn't instant:

- Process A takes time to write all data (e.g., 1 second).

2. Reads can happen mid-write:

- Process B might read **half-old + half-new data**.

Fix: Use a **mutex** so Process B waits until Process A finishes writing!

Fast Learner Activity:--

1.

Consider the following threads, T_1, T_2 , and T_3 executing on a single processor, synchronized using three binary semaphore variables, S_1, S_2 , and S_3 , operated upon using standard `wait()` and `signal()`. The threads can be context switched in any order and at any time.

T_1	T_2	T_3
<pre>while(true) { wait(S_3); print("C"); signal(S_2); }</pre>	<pre>while(true) { wait(S_1); print("B"); signal(S_3); }</pre>	<pre>while(true) { wait(S_2); print("A"); signal(S_1); }</pre>

Which initialization of the semaphores would print the sequence BCABCABC ...?

→

Initialization:

$S_1 = 1, S_2 = 0, S_3 = 0$

Execution Order:

- T_2 runs first (waits on $S_1=1$ → prints **B** → signals S_3).
- T_1 runs next (waits on $S_3=1$ → prints **C** → signals S_2).
- T_3 runs last (waits on $S_2=1$ → prints **A** → signals S_1).
- Repeat: **B** → **C** → **A** → **B** → **C** → **A**...

Why?

- $S_1=1$ allows T_2 (prints **B**) to start.
- $S_2=0$ and $S_3=0$ block T_3 and T_1 until signaled.
- Cycle: $T_2 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2 \dots$ produces **BCA** repeatedly.

2.

Consider the following two threads T_1 and T_2 that update two shared variables a and b . Assume that initially $a = b = 1$. Though context switching between threads can happen at any time, each statement of T_1 or T_2 is executed atomically without interruption.

T_1	T_2
$a = a + 1;$	$b = 2 * b;$
$b = b + 1;$	$a = 2 * a;$

Which one of the following options lists all the possible combinations of values of a and b after both T_1 and T_2 finish execution?

Possible final values of (a, b):

1. (3, 3)

- Occurs when T2 executes fully before T1:

$$b = 2*1 = 2 \rightarrow a = 2*1 = 2 \rightarrow a = 2+1 = 3 \rightarrow b = 2+1 = 3$$

2. (4, 3)

- Occurs in 4 interleavings (e.g., partial overlap):

$$b = 2*1 = 2 \rightarrow a = 1+1 = 2 \rightarrow a = 2*2 = 4 \rightarrow b = 2+1 = 3$$

3. (4, 4)

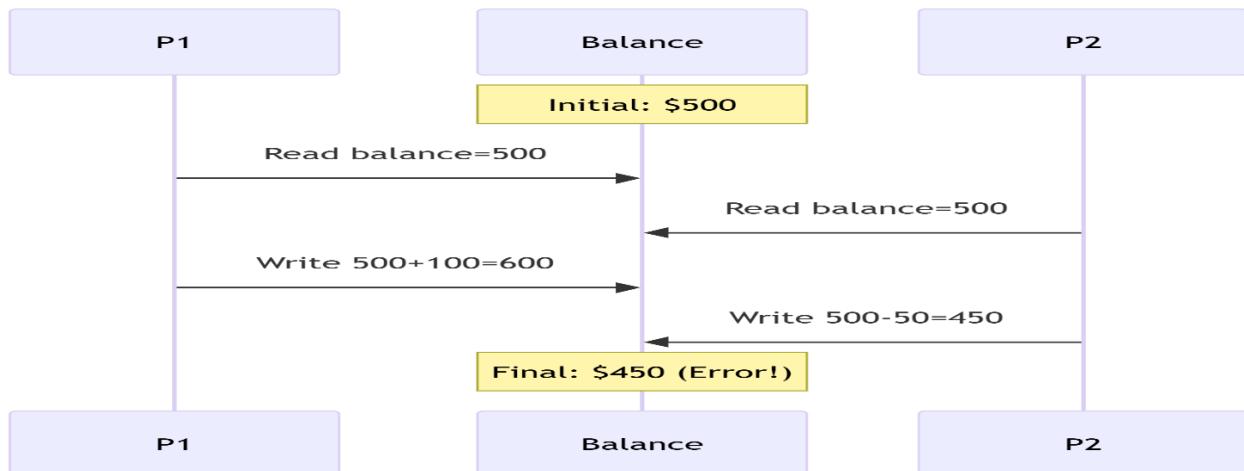
- Occurs when T1 executes fully before T2:

$$a = 1+1 = 2 \rightarrow b = 1+1 = 2 \rightarrow b = 2*2 = 4 \rightarrow a = 2*2 = 4$$

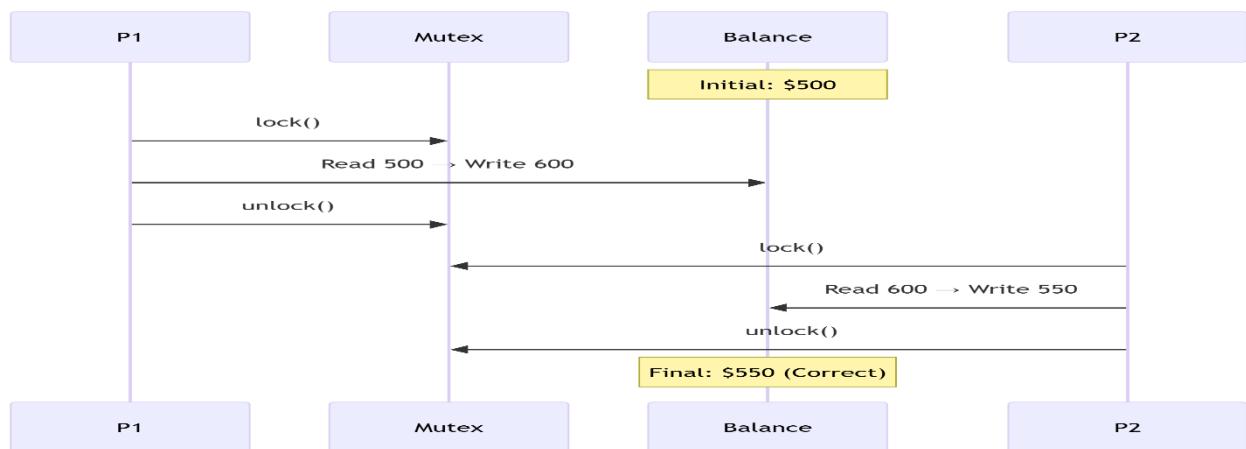
Impossible outcome: (3, 4)

- No execution order produces this combination.

3.No Synchronization:--



With Mutex:--



Based on the diagrams:

1. Why does the first diagram result in ticket=-1 and double booking?
2. How does the mutex in the second diagram prevent this?
3. What synchronization principle does this scenario demonstrate?

→ Answer Summary:

1. **Race Condition:** Both users see ticket_available simultaneously before either reserves.
2. **Mutex Serialization:**
 - User 1 locks → completes entire transaction → unlocks.
 - User 2 sees ticket=0 after lock.
3. **Principle: Mutual Exclusion** - Only one process accesses critical section (ticket reserve) at a time.

Without synchronization: Corrupt data (negative tickets).

With synchronization: Consistent state (no overbooking).